

IQMEETEQ WORKSHOP COMPANION

Your Workbook

PARTICIPANT PDF Tester

COHORT Tier 4 CEA Leadership — Day 5

GENERATED September 17, 2026

CONTENTS

Table of Contents

LEVEL 3 — AI BUILDER

The Idea
Build Your PRD
Closing

AI CHANGE LEADERSHIP

Three Ways to Put AI to Work
Consider the Work
Anatomy of Working Software
Build Your Mini-PRD
AI, ML, and LLMs
The Tools You Have
Many Ways In
Draft with Instructions and Role
The RICECO Framework
Real Work
Human in the Loop
Where It Earns Its Place
Pattern Matching and Hallucination
What AI Is / Is Not
Context
Create Your Context
LLM Peer Review
Power Follow-Ups
Cognitive Erosion
What Else Is Possible
Shadow AI
Data Safety
Red / Yellow / Green
The Change Message Framework
What You Know Now
Two Sentences

Three Concepts That Ride Along
Talk It Out

EXERCISE

The Idea

Two directions people arrive with: build something new, or replace something that doesn't fit.

WHAT DO YOU WANT TO BUILD OR REPLACE?

Sample answer for idea-build

WHO WILL USE IT?

Sample answer for idea-users

WHAT SHOULD IT DO?

No entry

WHAT'S FRUSTRATING ABOUT HOW YOU DO THIS TODAY?

No entry

WHAT WOULD MAKE IT FEEL LIKE A WIN?

No entry

Every build starts with one idea: something to make, or something to replace. Name yours before you prompt.

YOUR NOTES

No notes recorded

EXERCISE

Build Your PRD

Turn your idea into a prompt. Let Claude interview you.

PROMPT

You're going to help me plan an app I'm about to build with AI. I'm not a developer.

Interview me. Ask one question at a time and wait for my answer before asking the next. Keep your questions short and in plain English – no technical terms.

Work through these four things, in this order:

1. Who's allowed in, and how they get in
2. What people see and do when they open it
3. What rules it follows – what happens when someone does something
4. What it needs to remember between visits

If I don't know an answer, suggest something reasonable and move on. If my idea is getting too big, tell me and help me cut it down to the smallest version worth building.

When we're done, write me a short prompt I can paste into Replit. Write it as a person describing what they want – plain sentences, no bullet lists, no technical terms, no mention of databases or components. Six sentences or fewer.

Start by asking me what I want to build.

Open Claude (<https://claude.ai>)

Or write it yourself — four fields, same four questions.

WHO GETS IN, AND HOW?

No entry

WHAT DO PEOPLE SEE AND DO?

No entry

WHAT ARE THE RULES?

No entry

WHAT DOES IT NEED TO REMEMBER?

No entry

Building the app takes five minutes. The thinking you did to get here was the actual work.

THE TAKEAWAY

Consideration: One way to sharpen your build is to use an AI to help you think through what you want first, then bring that into Replit. A clearer starting point usually means fewer rounds. Some builds are simple enough to just prompt directly — this is a tool, not a rule.

YOUR NOTES

No notes recorded

EXERCISE

Closing

AN APP

Your idea, live on a URL.

A REPO

Somewhere your work survives.

A LOOP

Prompt. Look. Say more.

You can build software by describing it now. The tool got easier; the thinking is still yours.

Feedback (<https://example.com/feedback>)

YOUR NOTES

No notes recorded

REFERENCE

Three Ways to Put AI to Work

CHAT

You ask, it answers. Copilot and other chatbots. For drafting, summarizing, thinking out loud.

AUTOMATE

You set the steps, it runs them. n8n, Power Automate. For work that happens the same way every time.

DELEGATE

You give it a goal, it finds the path. A task, software, or a whole workflow — Replit, Lovable, Copilot Studio. For work where the steps are not known in advance.

On the left, you do the work. On the right, it does the work.

What does my division do by hand that it shouldn't?

YOUR NOTES

No notes recorded

EXERCISE

Consider the Work

GOAL

Something your division could use.

You are not going to build software today. You will be speaking it into existence. Describe it well enough that what comes back is what you meant. Same skill you use with IT, with a vendor, and with the person two doors down.

1 Look at your own division

Not the whole department. The work you are accountable for.

2 Name what is done by hand

Where the spreadsheet is. Who emails whom. Where it stalls.

3 Pick one thing

Small. A tracker, a checklist, a form that becomes a summary.

4 Write it in one sentence

What it should do that nobody has to do by hand anymore.

YOUR ONE THING

What should it do that nobody has to do by hand anymore?

No entry

HONEST GUESS

About 6 hours a week goes to reformatting status reports.

WHY

Nobody owns the template, so every team reinvents it.

YOUR NOTES

Free-form thoughts go here.

Second line.

REFERENCE

Anatomy of Working Software

You're opening a restaurant. How do people get in?

THE DOOR

Login. Who gets in, and what each role can see.

THE DINING ROOM

Front end. What people see and touch.

THE KITCHEN

Logic. Where the rules run.

THE FREEZER

Database. Where everything is stored.

Describe those four and you have written a Product Requirements Document.

YOUR NOTES

No notes recorded

EXERCISE

Build Your Mini-PRD

GOAL

One build per table.

1 Pick a scribe

Someone to type into the Workshop Companion for the group.

2 Document your requirements

Fill in the boxes below, then Copy and Submit. Anthony builds from what you submit.

TABLE NUMBER

Required.

4

TABLE NAME

Optional.

Night Owls

ONE SENTENCE

What should it do that nobody has to do by hand anymore?

A tool that turns meeting notes into a decision log.

LOGIN

Who gets in, and what each role can see.

No entry

FRONT END

What people see and touch.

Simple list of decisions with filters

LOGIC

Where the rules run.

No entry

DATABASE

Where everything is stored.

Decisions, owners, due dates

REQUIREMENTS AND CONSIDERATIONS

Anything the build has to respect.

No entry

SUBMITTED

Table number: 4

Table name: Night Owls

One sentence: A tool that turns meeting notes into a decision log.

Your builds

Links to the working prototypes will be added here during lunch.

YOUR NOTES

No notes recorded

REFERENCE

AI, ML, and LLMs

Artificial Intelligence is the broad field. Machine Learning is a subset of AI — it is already in your life: fraud detection at Visa, Face ID, Netflix recommendations, Outlook spam filtering, Google Maps routing. Large Language Models are a subset of Machine Learning — Claude, ChatGPT, Gemini — and they power tools like Copilot, NotebookLM, Notion AI, Canva AI, and Gamma.

LARGE

A vast knowledge base, trained on an enormous amount of text.

LANGUAGE

Natural, like you're typing or speaking to a person.

MODEL

The name of the tool itself: Claude, GPT, Gemini, and the rest.

The goal isn't to know everything. It's to know enough to lead, experiment, and decide.

YOUR NOTES

No notes recorded

REFERENCE

The Tools You Have

At work.

MICROSOFT COPILOT

General purpose. Drafting, summarizing, reorganizing, planning. Works across Outlook, Teams, Word, and SharePoint inside the Caltrans tenant. Best at language tasks on your own material. Weak at facts it was not given, numbers, anything it has to look up.

TRAFFIC MOBILITY INSIGHT

Purpose-built. Answers questions against a defined mobility dataset. Narrow, and accurate inside its lane. Best at the questions it was built for. Weak at anything outside the data it was pointed at.

ROAD SAFETY INSIGHT

Purpose-built. Same architecture, different data: crash and safety records. Best at pattern questions on safety data. Weak at general drafting. It is not a writing tool.

At home: Claude, ChatGPT, Gemini.

YOUR NOTES

No notes recorded

REFERENCE

Many Ways In

How are you using LLMs today? Every one of these is a language task. None requires the tool to know a fact it was not given.

DRAFT

Memos, emails, talking points. The blank page problem.

SYNTHESIZE

Five reports into one view. Twenty comments into three themes.

DISTILL

Forty pages into one. What survives is what matters.

BRAINSTORM

Options and angles you had not considered.

PLAN

A sequence. An agenda. The questions before a vendor call.

YOUR NOTES

No notes recorded

EXERCISE

Draft with Instructions and Role

GOAL

See what each part of a prompt changes.

Run these three in order, in a new Copilot chat each time. Notice what changes.

PROMPT

Write an email about a new policy requiring staff to complete training within 30 days.

PROMPT

You are a supportive supervisor who knows your team is stretched thin. Write an email about a new policy requiring staff to complete training within 30 days.

PROMPT

You are a construction foreman who runs a tight ship. Write an email about a new policy requiring staff to complete training within 30 days.

DEBRIEF

How did the role change it?

No entry

YOUR NOTES

No notes recorded

REFERENCE

The RICECO Framework

GOAL

Six ingredients for prompts that produce usable output the first time.

The RICECO Framework

LETTER	ELEMENT	QUESTION
R	Role	Who is the AI acting as?
I	Instruction	What exactly do you want it to do?
C	Context	What background information is needed?
E	Examples	What does good look like?
C	Constraints	What rules must it follow?
O	Output Format	How should the final result be formatted?

The 80% Shortcut

You don't need all six every time. For most daily tasks, I + C + C — Instruction, Context, Constraints — gets you 80% of the way there.

YOUR NOTES

No notes recorded

EXERCISE

Real Work

GOAL

Leave with one finished.

Your real queue. Not the important things — the recurring, language-heavy things. "A briefing memo due Thursday." "Forty public comments to summarize." The one you name is the one you work on. Real material. Real deadline.

DRAFT

Create something new — email, report, talking points, agenda

Ex: Memos, emails, talking points. The blank page problem.

No entry

SYNTHESIZE

Combine many things into one view

Ex: Five reports into one view. Twenty comments into three themes.

No entry

DISTILL

Shrink something long down to what matters

Ex: Forty pages into one. What survives is what matters.

No entry

BRAINSTORM

Generate options and angles

Ex: Options and angles you had not considered.

No entry

PLAN

Sequence the work

Ex: A rollout sequence. An agenda. The questions before a vendor call.

No entry

YOUR NOTES

No notes recorded

REFERENCE

Human in the Loop

The fix isn't better AI. It's a human in the loop. Beginning: you direct. Middle: AI processes. End: you edit and oversee it.

The trap

These systems perform trust. They don't earn it. The better they get at the performance, the harder it is to remember. Your job: understand it. Don't trust it.

YOUR NOTES

No notes recorded

REFERENCE

Where It Earns Its Place

Which of your real tasks suit the tools you actually have? Recurring, language-heavy work. Not anything the number has to be right on.

An agent is different. You hand it a goal. It decides what files to create, writes them, hits errors, reads the errors, fixes them, and keeps going until something ran. Their words went in. An agent did the rest.

YOUR NOTES

No notes recorded

Pattern Matching and Hallucination

An LLM predicts the next word based on massive training data. It's not retrieving a fixed answer; it's generating a likely sequence. When it lands on truth, we call it an answer. When it lands on false, we call it a hallucination. Same process either way. It doesn't know it's right. It doesn't know it's wrong. It's always guessing.

More context means better predictions. "Mary ___" could be anything. "Mary had a little ___" is almost certainly one word. Vague prompts get vague outputs. Your context is the lever.

You: That statistic you cited — where did it come from? AI: It's from the 2023 National Workforce Study by the Bureau of Labor Statistics. You: I can't find that study. Does it exist? AI: I filled in a pattern that I thought belonged there.

YOUR NOTES

No notes recorded

REFERENCE

What AI Is / Is Not

GOAL

Pattern matching, not thinking. The same process produces correct answers and hallucinations.

Core Concept

AI is pattern matching, not thinking. The same process that produces correct answers also produces hallucinations. It doesn't know the difference.

What This Means for Your Work

AI can draft, brainstorm, summarize, and restructure. It cannot verify facts, exercise professional judgment, or understand the human stakes of your decisions. Every output needs a human checkpoint before it reaches a client, a colleague, or a decision-maker.

YOUR NOTES

No notes recorded

Context

ONE CHAT

Everything typed so far. That is what it works from. A fouled chat stays fouled — the bad assumption is in the context.

WHAT IT CAN REACH

Files, mail, documents. Some tools read what you point them at. Whether yours does is a configuration question, not a capability one.

ACROSS SESSIONS

What it remembers. Claude and GPT carry memory and projects. Copilot chat, as you have it, starts fresh each time.

TO AN AGENT

What you hand it, plus what it fetches. The leap that makes agents useful and risky at once.

Bad context is why more prompting stops helping. Know which layer the problem is in.

DRIFT

The topic wandered. Twenty exchanges in, the chat is somewhere else. Start a new one.

FOULED

A bad assumption got baked in. Every answer now leans into it. Edit the earlier prompt, or abandon the chat.

BAD REPLIES

Same wrong thing, again. More prompting will not fix it. The context is the problem. Sunset it.

YOUR NOTES

No notes recorded

EXERCISE

Create Your Context

GOAL

What do you want it to remember about you?

1 Pick a task where the tool would need to know you

Where a generic answer wouldn't work, because of your audience, your division, or how it has to sound.

2 Fill in the boxes that would have the most impact

Role, acronyms, formatting. Skip the rest.

3 Copy, paste into a new Copilot chat, then type the task

How did your context document influence it?

ROLE

What should AI know about your role?

Your title or function and where you sit in the organization.

No entry

WORK

What should AI know about the work you're responsible for?

Recurring responsibilities that affect the help you need.

No entry

TERMS

What terminology or acronyms should AI understand without explanation?

Only terms you use regularly.

No entry

RULES

What standing rules or constraints should AI always account for?

Requirements that apply across most of your work.

No entry

tone

What should AI know about your preferred tone?

How responses should generally sound.

No entry

FORMAT

How do you prefer information to be organized?

BLUF, bullets, sections, tables, or narrative.

No entry

DEPTH

How much detail do you usually want?

Your normal level of depth or concision.

No entry

DELIVERABLES

What kinds of deliverables do you commonly want AI to produce?

Emails, memos, briefings, talking points, or analyses.

No entry

AVOID

What should AI avoid doing?

Recurring response habits that make the output less useful.

No entry

YOUR NOTES

No notes recorded

EXERCISE

LLM Peer Review

GOAL

Use a second AI to check the first one's work. Build the verification habit.

The Exercise

You ran the same task through two different AI tools and compared their outputs. Then you used one model to critique the other's work.

The Critique Scaffold

"Here are the instructions I gave to [Model A]. Below is the response. Critique only. Do not redraft."

Why This Matters

AI models have different strengths and blind spots. Using a second model as a reviewer catches errors that a single model misses — especially factual claims, logical gaps, and tone issues.

YOUR NOTES

No notes recorded

Power Follow-Ups

GOAL

Nine moves to refine, pressure-test, and reshape AI output until it's actually ready to use.

Nine Moves to Refine AI Output

1. **Go Deeper** — "Expand on point 3 with specific examples."
2. **Change Format** — "Rewrite this as a table / email / FAQ / one-pager."
3. **Shift Audience** — "Rewrite for [board members / new staff / the public]."
4. **Challenge It** — "What are the strongest counterarguments to this?"
5. **Simplify** — "A smart 8th grader should understand this. Rewrite."
6. **Add Constraints** — "Now do it in under 200 words / without jargon / in Spanish."
7. **Verify** — "What sources support these claims? Flag anything you're uncertain about."
8. **Compare** — "How does this compare to [alternative approach]?"
9. **Pressure-Test** — "What's missing? What would a skeptic say?"

These work with any AI tool, on any task. Use them after your first prompt to push the output from "okay" to "actually useful."

YOUR NOTES

No notes recorded

Cognitive Erosion

It's not laziness. It's misplaced confidence. Confidence in the tool lowers critical thinking; confidence in yourself raises it. Effort drops most at the judgment end — evaluation, synthesis, analysis. Tool first, thinking second: the order matters.

The failure is invisible until a decision goes wrong. Gradual drift: session one you edit, session three you copy-paste. Output homogenizes — the memos start sounding alike. "It wrote it" replaces "I wrote it." Skill loss shows up only when the tool is gone. Two or three people trusted the polish before it reached you.

ASK WHAT TO CONSIDER

Ask what to consider and why, before you ask for the answer.

THINK FIRST

Draft the skeleton yourself, then use the tool.

MAKE IT ARGUE

Have it argue the other side of your draft. "What would an expert check here that I haven't?"

HAVE IT QUIZ YOU

On your own document, before you send it.

What a leader does

Ask where it came from. Keep one hard thing by hand. Reward the catch. Model it yourself. Prevention is structural: the tool is designed to make carelessness feel like competence.

YOUR NOTES

No notes recorded

REFERENCE

What Else Is Possible

Beyond Copilot, Traffic Mobility Insight, and Road Safety Insight. Remember the three ways: Chat, Automate, Delegate.

Automate — the path is fixed, AI can live inside it

- 1 Trigger**
Something happens.
- 2 AI step**
Summarize, classify, extract.
- 3 Action**
Update, route, create.
- 4 Output**
Notify, store, send.

The workflow does not decide what happens next. You already did.

Delegate — the path is discovered while the work is happening. Example: "Prepare Monday's executive brief." The agent pulls the calendar, project data, email updates, and policy news, asks what's missing, searches an alternate source, synthesizes, drafts. The path changes when the evidence changes.

Some of the people you lead have used these tools. They know the gap.

YOUR NOTES

No notes recorded

Shadow AI

Shadow AI is evidence. It tells leaders where capability, policy, and user need are out of alignment.

You did not choose it. Copilot arrived in an update — every M365 tenant got it. Gemini arrived in a browser patch that IT approved as a security fix. Connectors arrive through your own tools: an approved tool with an admin panel connects to anything. The question is not whether you have shadow AI. It is whether you can see it.

Read the signal — the workaround is telling you something

CAPABILITY

The approved tool cannot do the job. People reach outside when the sanctioned option stops short.

POLICY

The rules do not match the work. A blanket restriction collides with a legitimate operational need.

ENABLEMENT

People do not know the sanctioned path. The capability may exist; they just don't know how to reach it.

Before you punish the workaround, diagnose the gap.

"WHO IS USING UNAPPROVED AI?"

Finds a person.

"WHAT ARE THEY TRYING TO ACCOMPLISH THAT THEY CANNOT ACCOMPLISH ANOTHER WAY?"

Finds the gap.

SURVEILLANCE

Catch it. Stop it. People learn what not to tell you. The behavior gets harder to see, not less common.

SUNLIGHT

See it. Understand it. Ask what the tool is doing for them. Visibility gives you something you can actually govern.

Make the sanctioned path win. People choose the approved path when it gives them more, not merely when policy tells them to: approved data, shared prompts, internal knowledge, better integrations, reusable workflows. Compliance is not the incentive. Capability is.

Trust but verify

Enable, train, make the path known — that sets the direction and gets people to tell you what they use. Then check at machine speed. AI runs at machine speed, so the check on it has to.

HONEST GUESS

What is your division already using that you cannot see?

Not what they may use. What they use. Personal ChatGPT. Claude on a phone.

No entry

WHY

What need is that tool meeting that the sanctioned path does not?

No entry

MONDAY

Sunlight or surveillance? Say which, and name the first move.

No entry

YOUR NOTES

No notes recorded

Data Safety

Before we talk about AI safety, let's talk about data safety. Optimize entirely for efficiency and everything is wide open. Optimize entirely for safety and we'd never send an email. Every organization lives somewhere in the middle.

"Is this safe?" is not a question that has an answer. Safe compared to what? For which material? Against which alternative you already use?

It already happened. In 2024, researchers extracted over 2,700 real credentials — API keys, passwords, access tokens — from GitHub's AI, which had learned them during training and suggested them to other developers. When model improvement is on, your inputs may be used to improve the product. Turning model improvement off is a controllable setting, and it should be checked before any sensitive work.

How your tools handle your data

OUTLOOK & SHAREPOINT

Stored on Microsoft's servers. Encrypted. Not used to build products for others. Kept until you delete it.

PAYROLL / HR

Stored on the vendor's servers. Encrypted. Not used to build products for others. Kept 4+ years by federal law.

COPILOT, CALTRANS TENANT

Stored in Microsoft's government cloud. Encrypted. Not used to build products for others, by contract. Kept per the state agreement.

CONSUMER AI

Stored on the AI company's servers. Encrypted. Used to build products for others only if model improvement is on. Kept until you delete it.

The security question is not "is AI safe?" It's "what's the data, what are the settings, and what does my organization allow?"

For a staffer, safety is a toggle. For you, safety is a policy: which tools your people may use, and how you find out when they use something else.

YOUR NOTES

No notes recorded

EXERCISE

Red / Yellow / Green

GOAL

Build shared judgment about what's safe, context-dependent, and risky to do with AI at work.

The Exercise

Your team sorted AI use cases into three categories based on risk level and built shared judgment about what's appropriate in your work context.

YOUR NOTES

No notes recorded

REFERENCE

The Change Message Framework

You already wrote it. Now say it in order.

- 1 Forces**
What changed outside.
- 2 Practices**
How we do it today.
- 3 Consequences**
What that costs now.
- 4 A stronger path**
What we saw.
- 5 Impact**
What is different in 90 days.

Five sentences. Repeated in the all-hands, the hallway, and the one-on-one until it sticks.

SKIP THE COST

"Here is a better way." Lands as a preference.
Preferences lose to the status quo.

NAME THE COST FIRST

"Here is what today costs. Here is the better way."
Lands as a recovery. The same words, in that order,
move people.

The brain registers a gain only against a loss. Steps 2–3 are the loss. Steps 4–5 are the gain. Skip either half and the message stops working.

YOUR NOTES

No notes recorded

EXERCISE

What You Know Now

GOAL

One sentence each. Eight minutes.

OUTSIDE

What changed around your division that nobody in it chose?

No entry

TODAY

How does your division do that work right now? Say it without judgment.

No entry

COST

Knowing what you know now, what is that costing?

No entry

SAW IT

What did you see built today that changes that?

No entry

90 DAYS

What is different in 90 days if it works?

No entry

Hand it to the person next to you. They read your five sentences out loud, in order. Which sentence is weakest? Which one would your division argue with?

YOUR NOTES

No notes recorded

EXERCISE

Two Sentences

GOAL

The ones you would actually say.

Rewrite rows 1 and 2 as speech. Who hears it: a direct report who did not ask. Competent, busy, quietly skeptical. Then say it to your partner. They answer one question only: "What did you just hear me asking you to do?" If the answer is "nothing yet," the opening worked.

SENTENCE ONE

What changed outside your division.

No entry

SENTENCE TWO

How your division does that work today — said so they nod, not flinch.

No entry

YOUR NOTES

No notes recorded

REFERENCE

Three Concepts That Ride Along

HUMAN IN THE LOOP

Verify before it circulates. Facts, math, sources.

SUNLIGHT, NOT SURVEILLANCE

"I want to know what you already use." Not to stop you. To understand what it does for you.

WHERE IT EARNS ITS PLACE

Recurring, language-heavy. Not anything the number has to be right on.

Three concepts, no more. They repeat until your division can say them back.

YOUR NOTES

No notes recorded

EXERCISE

Talk It Out

WHAT CLICKED

One insight, technique, or moment that stuck.

No entry

WHO NEEDS TO HEAR IT

A peer CEA who was not in this room. What would you tell them first?

No entry

Tell us how we did — two minutes in the Feedback tab, now, while it is fresh. Tomorrow: Day 6, Succession Planning and Leadership Legacy. Same room. Bring the two sentences you wrote today.

YOUR NOTES

No notes recorded